

LSTM FOR TEXT CLASSIFICATION

In [1]:

```
import pandas as pd
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from keras.models import Model
from keras.layers import LSTM, Activation, Dense, Dropout, Input, Embedding
from keras.optimizers import RMSprop
from keras.preprocessing.text import Tokenizer
from keras.preprocessing import sequence
from keras.utils import to_categorical
from keras.utils import pad_sequences
from keras.callbacks import EarlyStopping
```

In [2]:

Load the data

In [2]:

```
df = pd.read_csv('/content/spam.csv', delimiter=',', encoding='latin-1')
df.head()
```

Out[2]:

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only	NaN	NaN	NaN
		...			
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	Udunsaysoearlyhor...Ucalreadythensay...	NaN	NaN	NaN
4	ham	NahIdon'tthinkhegoestousf,helivesaro...	NaN	NaN	NaN

In [4]:

Drop unnecessary columns

In [4]:

```
df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'], axis=1, in place=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'> Range
Index: 5572 entries, 0 to 5571
Data columns (total 2
columns):
```

#	Column	Non-Null Count	Dtype
0	v1	5572 non-null	object
1	v2	5572 non-null	object

In [5]:

Create input and output vectors and process the labels

In [1]:

```
X = df.v2 Y =  
df.v1  
le = LabelEncoder()  
Y = le.fit_transform(Y) Y =  
Y.reshape(-1,1)
```

In [6]:

split the data set for training and test.

In [2]:

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.15)
```

In [7]:

Process the data

In [8]:

```
max_words = 1000  
max_len = 150  
tok = Tokenizer(num_words=max_words)  
tok.fit_on_texts(X_train)  
sequences = tok.texts_to_sequences(X_train)  
sequences_matrix = tf.keras.utils.pad_sequences(sequences,maxlen=max_len)
```

In [9]:

Define the model

In [10]:

```
def RNN():  
    inputs = Input(name='inputs',shape=[max_len])  
    layer = Embedding(max_words,50,input_length=max_len)(inputs)  
    layer = LSTM(64)(layer)  
    layer = Dense(256,name='FC1')(layer)  
    layer = Activation('relu')(layer)  
    layer = Dropout(0.5)(layer)  
    layer = Dense(1,name='out_layer')(layer)  
    layer = Activation('sigmoid')(layer)  
    model = Model(inputs=inputs,outputs=layer)  
    return model
```

In [11]:

Call the function and compile the model

In [11]:

```
model = RNN()  
model.summary()  
model.compile(loss='binary_crossentropy',optimizer=RMS prop(),metrics=['accuracy'])
```

Model: "model_1"

Layer (type)

Output Shape

Param #

inputs (Input Layer)	[(None, 150)]	0
embedding_1 (Embedding)	(None, 150, 50)	50000
lstm_1 (LSTM)	(None, 64)	29440
FC1 (Dense)	(None, 256)	16640
activation_2 (Activation)	(None, 256)	0
dropout_1 (Dropout)	(None, 256)	0
out_layer (Dense)	(None, 1)	257
activation_3 (Activation)	(None, 1)	0

Total params: 96,337
Trainable params: 96,337
Non-trainable params: 0

In [12]:

Fit the model

In [1]:

```
model.fit(sequences_matrix, Y_train, batch_size=128, epochs=10,
          validation_split=0.2, callbacks=[EarlyStopping(monitor='val_loss', min_delta=0.001)])
```

Epoch 1/10
30/30 [=====] - 10s 267ms/step - loss: 0.3345 - accuracy:0.8730
- val_loss: 0.1491 - val_accuracy:0.9462 Epoch 2/10
30/30 [=====] - 8s 251ms/step - loss: 0.0887 - accuracy: 0.9794
- val_loss: 0.0625 - val_accuracy:0.9821 Out[54]:

<keras.callbacks.History at 0x7f0a5c167750>

In [13]:

Process the test data

In [1]:

```
test_sequences = tok.texts_to_sequences(X_test)
test_sequences_matrix = tf.keras.utils.pad_sequences(test_sequences, maxlen=max_len)
```

In [14]:

Evaluate the model with the test

In [1]:

```
accr = model.evaluate(test_sequences_matrix, Y_test)
```

27/27 [=====] - 1s 21ms/step - loss: 0.0643 - accuracy: 0.9797

In [15]:

```
print("Test set\n      Loss: {:.3f}\n      Accuracy: {:.3f}'.format(incr[0],incr[1])) Test set
      Loss: 0.064
      Accuracy: 0.980
```