

Assignment 3: Build CNN Model for Classification Of flowers.

```
Import splitfolders
import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.preprocessing import image
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
from tensorflow.keras.models import load_model
from tensorflow.keras.layers import
Dense,Convolution2D,MaxPooling2D,Flatten
from tensorflow.keras.applications.resnet50 import preprocess_input,
decode_predictions
from tensorflow.keras.preprocessing import image
import matplotlib.pyplot as plt
```

2. Image Augmentation

In [2]:

```
train_datagen =
ImageDataGenerator(rescale=1./255, zoom_range=0.2, horizontal_flip=True, ver
tical_flip=False)
```

In [3]:

```
test_datagen = ImageDataGenerator(rescale=1./255)
```

In [4]:

```
input_folder = 'C:\\Users\\manok\\Documents\\Sem_7\\HX5001-
HX6001\\Assignment\\Assignment_3\\flowers'
```

In [5]:

```
splitfolders.ratio(input_folder,output="C:\\Users\\manok\\Documents\\Sem_7\\
HX5001-
HX6001\\Assignment\\Assignment_3\\flowersdataset",ratio=(.8,0,.2),group_pre
fix=None)
```

```
Copying files: 4317 files [00:26, 166.01 files/s]
```

In [6]:

```
x_train=train_datagen.flow_from_directory(r"C:\Users\manok\Documents\Sem
_7\HX5001-
HX6001\Assignment\Assignment_3\flowersdataset\train",target_size=(64,64),cl
ass_mode='categorical',batch_size=24)
```

Found 3452 images belonging to 5 classes.

In [7]:

```
x_test=test_datagen.flow_from_directory(r"C:\Users\manok\Documents\Sem_7\HX5001-HX6001\Assignment\Assignment_3\flowersdataset\test",target_size=(64,64),class_mode='categorical',batch_size=24)
```

Found 865 images belonging to 5 classes.

In [8]:

```
x_train.class_indices
```

Out[8]:

```
{'daisy': 0, 'dandelion': 1, 'rose': 2, 'sunflower': 3, 'tulip': 4}
```

3. Create Model

In [9]:

```
model=Sequential()
```

4. Add Layers

4.1. Convolution Layer

In [10]:

```
model.add(Convolution2D(32,(3,3),input_shape=(64,64,3),activation='relu'))
```

4.2. MaxPooling Layer

In [11]:

```
model.add(MaxPooling2D(pool_size=(2,2)))
```

4.3. Flatten Layer

In [12]:

```
model.add(Flatten())
```

4.4. Dense Layer

In [13]:

```
model.add(Dense(300,activation='relu'))
model.add(Dense(150,activation='relu'))
```

In [14]:

```
model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
flatten (Flatten)	(None, 30752)	0
dense (Dense)	(None, 300)	9225900
dense_1 (Dense)	(None, 150)	45150
=====		
Total params: 9,271,946		
Trainable params: 9,271,946		
Non-trainable params: 0		

4.5. Output Layer

```

In [15]:
model.add(Dense(5,activation='softmax'))

In [16]:
model.summary()
Model: "sequential"

```

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 62, 62, 32)	896
max_pooling2d (MaxPooling2D)	(None, 31, 31, 32)	0
flatten (Flatten)	(None, 30752)	0

dense (Dense)	(None, 300)	9225900
dense_1 (Dense)	(None, 150)	45150
dense_2 (Dense)	(None, 5)	755

```
=====
=====
```

```
Total params: 9,272,701
Trainable params: 9,272,701
Non-trainable params: 0
```

```
—
```

5. Compile The Model

```
In [17]:
model.compile(loss='categorical_crossentropy',optimizer='adam',metrics=['accuracy'])
len(x_train)
```

```
Out[17]:
144
```

6. Fit The Model

```
In [18]:
epo=20
history =
model.fit(x_train,steps_per_epoch=len(x_train),validation_data=x_test,validation_steps=len(x_test),epochs=epo)

Epoch 1/20
144/144 [=====] - 81s 552ms/step - loss: 1.4602 - accuracy: 0.4429 - val_loss: 1.2587 - val_accuracy: 0.5052
Epoch 2/20
144/144 [=====] - 38s 264ms/step - loss: 1.0639 - accuracy: 0.5773 - val_loss: 1.1310 - val_accuracy: 0.5399
Epoch 3/20
144/144 [=====] - 38s 263ms/step - loss: 0.9872 - accuracy: 0.6066 - val_loss: 1.0271 - val_accuracy: 0.5861
Epoch 4/20
144/144 [=====] - 38s 263ms/step - loss: 0.9298 - accuracy: 0.6251 - val_loss: 1.0208 - val_accuracy: 0.6266
Epoch 5/20
```

144/144 [=====] - 38s 264ms/step - loss: 0.8497 - accuracy: 0.6651 - val_loss: 0.9911 - val_accuracy: 0.6428
Epoch 6/20
144/144 [=====] - 38s 263ms/step - loss: 0.8255 - accuracy: 0.6727 - val_loss: 1.1223 - val_accuracy: 0.6023
Epoch 7/20
144/144 [=====] - 36s 253ms/step - loss: 0.7639 - accuracy: 0.7048 - val_loss: 1.0702 - val_accuracy: 0.6243
Epoch 8/20
144/144 [=====] - 37s 254ms/step - loss: 0.7179 - accuracy: 0.7170 - val_loss: 1.1313 - val_accuracy: 0.5873
Epoch 9/20
144/144 [=====] - 37s 254ms/step - loss: 0.6676 - accuracy: 0.7352 - val_loss: 0.9532 - val_accuracy: 0.6647
Epoch 10/20
144/144 [=====] - 36s 252ms/step - loss: 0.6323 - accuracy: 0.7567 - val_loss: 0.9810 - val_accuracy: 0.6532
Epoch 11/20
144/144 [=====] - 37s 253ms/step - loss: 0.6231 - accuracy: 0.7590 - val_loss: 1.0481 - val_accuracy: 0.6439
Epoch 12/20
144/144 [=====] - 36s 254ms/step - loss: 0.5839 - accuracy: 0.7775 - val_loss: 1.0738 - val_accuracy: 0.6821
Epoch 13/20
144/144 [=====] - 36s 253ms/step - loss: 0.5251 - accuracy: 0.8097 - val_loss: 0.9613 - val_accuracy: 0.6682
Epoch 14/20
144/144 [=====] - 35s 245ms/step - loss: 0.4838 - accuracy: 0.8143 - val_loss: 1.0360 - val_accuracy: 0.6682
Epoch 15/20
144/144 [=====] - 96s 667ms/step - loss: 0.4308 - accuracy: 0.8433 - val_loss: 1.1060 - val_accuracy: 0.6647
Epoch 16/20
144/144 [=====] - 26s 180ms/step - loss: 0.4230 - accuracy: 0.8491 - val_loss: 1.2172 - val_accuracy: 0.6590
Epoch 17/20
144/144 [=====] - 28s 192ms/step - loss: 0.4122 - accuracy: 0.8517 - val_loss: 1.0914 - val_accuracy: 0.6694
Epoch 18/20
144/144 [=====] - 29s 199ms/step - loss: 0.3877 - accuracy: 0.8644 - val_loss: 1.4504 - val_accuracy: 0.5988
Epoch 19/20

```

144/144 [=====] - 39s 271ms/step - loss: 0.
3670 - accuracy: 0.8653 - val_loss: 1.2226 - val_accuracy: 0.6428
Epoch 20/20
144/144 [=====] - 39s 272ms/step - loss: 0.
3131 - accuracy: 0.8853 - val_loss: 1.2005 - val_accuracy: 0.6798


```

In [19]:

```

epochs_range = range(epo)
plt.figure(figsize=(8, 8))
plt.plot(epochs_range, history.history['accuracy'], label='Training Accuracy')
plt.plot(epochs_range, history.history['val_accuracy'], label='Validation Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy')
plt.show()


```



```

plt.figure(figsize=(8, 8))
plt.plot(epochs_range, history.history['loss'], label='Training Loss')
plt.plot(epochs_range, history.history['val_loss'], label='Validation Loss')
plt.legend()
plt.title('Training and Validation Loss')
plt.show()

```



7. Save the Model

In [21]:

```
model.save('flowers.h5')
```

8. Test the Model

In [22]:

```
img=image.load_img(r"C:\Users\manok\Documents\Sem_7\HX5001-
HX6001\Assignment\Assignment_3\flowersdataset\test\daisy\3706420943_66f
3214862_n.jpg",target_size=(64,64))
x=image.img_to_array(img)
x=np.expand_dims(x,axis=0)
y=np.argmax(model.predict(x),axis=1)
x_train.class_indices
index=['daisy','dandellion','rose','sunflower','tulip']
index[y[0]]
```

1/1 [=====] - 1s 661ms/step

Out[22]:

'daisy'

In [23]:

```
img_url =
"https://storage.googleapis.com/download.tensorflow.org/example_images/592p
x-Red_sunflower.jpg"
img_path = tf.keras.utils.get_file('Red_sunflower', origin=img_url)
```

```
img = image.load_img(img_path, target_size=(224, 224))
img_array = image.img_to_array(img)
img_batch = np.expand_dims(img_array, axis=0)
```

```
img_preprocessed = preprocess_input(img_batch)
model = tf.keras.applications.resnet50.ResNet50()
prediction = model.predict(img_preprocessed)

print(decode_predictions(prediction, top=3)[0])

score = tf.nn.softmax(prediction[0])
1/1 [=====] - 2s 2s/step
[('n11939491', 'daisy', 0.57757616), ('n02206856', 'bee', 0.249383), ('n03991062', 'pot', 0.01181932)]
```